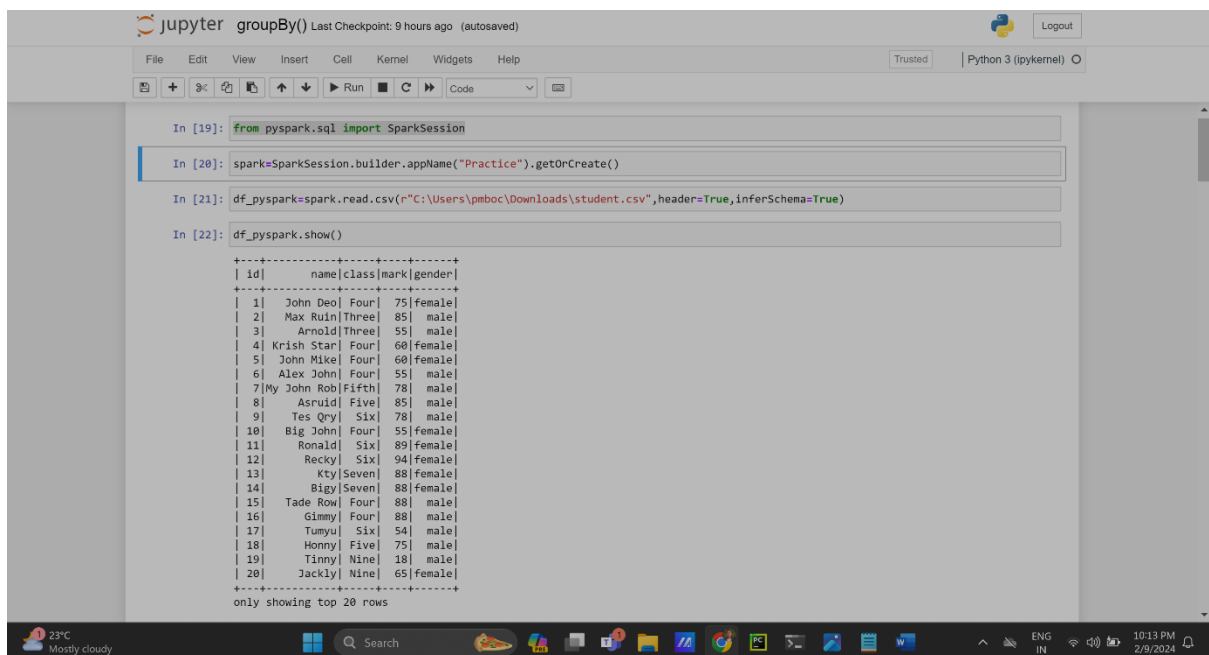


Name: Pradip Bochare

GroupBy and Aggregate function

Similar to SQL GROUP BY clause, PySpark groupBy() function is used to collect the identical data into groups on DataFrame and perform count, sum, avg, min, and max functions on the grouped data.



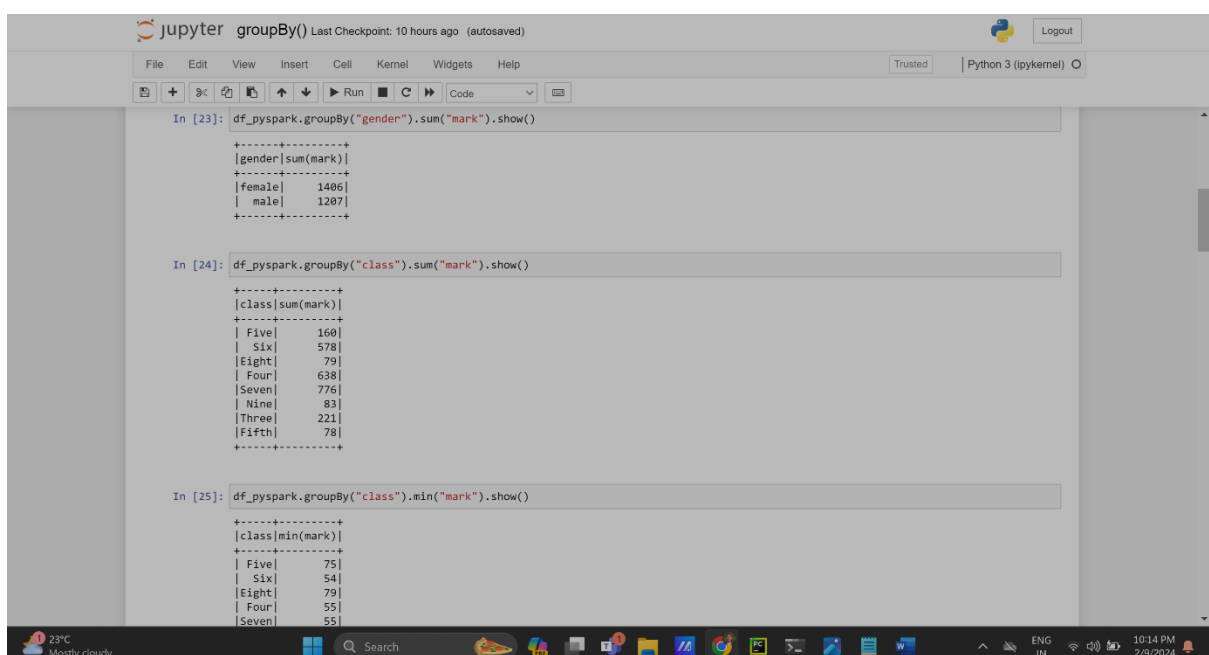
```
In [19]: from pyspark.sql import SparkSession

In [20]: spark=SparkSession.builder.appName("Practice").getOrCreate()

In [21]: df_pyspark=spark.read.csv(r"C:\Users\pmboch\Downloads\student.csv",header=True,inferSchema=True)

In [22]: df_pyspark.show()

+-----+-----+-----+-----+
| id | name | class | mark | gender |
+-----+-----+-----+-----+
| 1 | John Deo | Four | 75 | female |
| 2 | Max Ruin | Three | 85 | male |
| 3 | Arnold | Three | 55 | male |
| 4 | Krish Star | Four | 60 | female |
| 5 | John Mike | Four | 60 | female |
| 6 | Alex John | Four | 55 | male |
| 7 | My John Rob | Fifth | 78 | male |
| 8 | Asrud | Five | 85 | male |
| 9 | Tes Qryl | Six | 78 | male |
| 10 | Big John | Four | 55 | female |
| 11 | Ronald | Six | 89 | female |
| 12 | Recky | Six | 94 | female |
| 13 | Kty | Seven | 88 | female |
| 14 | Bigy | Seven | 88 | female |
| 15 | Tade Row | Four | 88 | male |
| 16 | Gimmy | Four | 88 | male |
| 17 | Tumyu | Six | 54 | male |
| 18 | Honny | Five | 75 | male |
| 19 | Tinny | Nine | 18 | male |
| 20 | Jackly | Nine | 65 | female |
+-----+-----+-----+-----+
only showing top 20 rows
```



```
In [23]: df_pyspark.groupBy("gender").sum("mark").show()

+-----+-----+
| gender | sum(mark) |
+-----+-----+
| female | 1406 |
| male | 1207 |
+-----+-----+

In [24]: df_pyspark.groupBy("class").sum("mark").show()

+-----+-----+
| class | sum(mark) |
+-----+-----+
| Five | 160 |
| Six | 578 |
| Eight | 79 |
| Four | 638 |
| Seven | 776 |
| Nine | 83 |
| Three | 221 |
| Fifth | 78 |
+-----+-----+

In [25]: df_pyspark.groupBy("class").min("mark").show()

+-----+-----+
| class | min(mark) |
+-----+-----+
| Five | 75 |
| Six | 54 |
| Eight | 79 |
| Four | 55 |
| Seven | 55 |
+-----+-----+
```

Jupyter groupBy() Last Checkpoint: 10 hours ago (autosaved) Python 3 (pykernel)

File Edit View Insert Cell Kernel Widgets Help Trusted

```
In [26]: df_pyspark.groupBy("class").max("mark").show()
```

class	max(mark)
Five	85
Six	96
Eight	79
Four	88
Seven	90
Nine	65
Three	85
Fifth	78

```
In [27]: df_pyspark.groupBy("class").avg("mark").show()
```

class	avg(mark)
Five	80.0
Six	82.57142857142857
Eight	79.0
Four	70.88888888888889
Seven	77.6
Nine	41.5
Three	73.66666666666667
Fifth	78.0

```
In [28]: df_pyspark.groupBy("class").mean("mark").show()
```

23°C Mostly cloudy 10:14 PM 2/9/2024

Jupyter groupBy() Last Checkpoint: 10 hours ago (autosaved) Python 3 (pykernel)

File Edit View Insert Cell Kernel Widgets Help Trusted

```
In [28]: df_pyspark.groupBy("class").mean("mark").show()
```

class	avg(mark)
Five	80.0
Six	82.57142857142857
Eight	79.0
Four	70.88888888888889
Seven	77.6
Nine	41.5
Three	73.66666666666667
Fifth	78.0

```
In [29]: df_pyspark.groupBy("class").count().show()
```

class	count
Five	2
Six	7
Eight	1
Four	9
Seven	10
Nine	2
Three	3
Fifth	1

```
In [30]: df_pyspark.groupBy("class").pivot("name").sum("mark").show()
```

23°C Mostly cloudy 10:15 PM 2/9/2024

🔗 orderBy() and sort() in Pyspark DataFrame

The screenshot shows a Jupyter Notebook interface with the following content:

At the top, the Jupyter logo and the text "Last Checkpoint: 10 hours ago (autosaved)" are visible. The interface includes a menu bar (File, Edit, View, Insert, Cell, Kernel, Widgets, Help) and a toolbar with icons for file operations, running, and code execution. The Python 3 (pykernel) environment is selected.

The first code cell (In [31]:) contains the following code:

```
df_pyspark.sort("mark").show()
```

The output of this cell is a table showing the top 20 rows of the DataFrame sorted by 'mark' in ascending order. The table has columns: id, name, class, mark, and gender. The rows are as follows:

id	name	class	mark	gender
19	Tinny	Nine	18	male
17	Tumyu	Six	54	male
3	Arnold	Three	55	male
6	Alex John	Four	55	male
10	Big John	Four	55	female
22	Reggid	Seven	55	female
29	Tess Played	Seven	55	male
4	Krish Star	Four	60	female
5	John Mike	Four	60	female
20	Jackly	Nine	65	female
21	Babby John	Four	69	female
34	Gain Toe	Seven	69	male
1	John Deo	Four	75	female
18	Honny	Five	75	male
7	My John Rob	Five	78	male
9	Tes Qryl	Six	78	male
24	Tiddy Now	Seven	78	male
23	Herod	Eight	79	male
26	Crelea	Seven	79	male
30	Reppy Red	Six	79	female

Below the table, the text "only showing top 20 rows" is displayed.

The second code cell (In [32]:) contains the following code:

```
df_pyspark.sort(df_pyspark["mark"].desc()).show()
```

The output of this cell is a table showing the top 20 rows of the DataFrame sorted by 'mark' in descending order. The table has columns: id, name, class, mark, and gender. The rows are as follows:

id	name	class	mark	gender
----	------	-------	------	--------

The table is partially visible, showing the header and the first row (id: 1, name: name, class: class, mark: mark, gender: gender).

Jupyter groupBy() Last Checkpoint: 10 hours ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (pykernel)

```
In [32]: df_pyspark.sort(df_pyspark["mark"].desc()).show()
```

id	name	class	mark	gender
33	Kenn Rein	Six	96	Female
12	Recky	Six	94	Female
32	Binn Rott	Seven	90	Female
11	Ronald	Six	89	Female
13	Kty	Seven	88	Female
14	Bigy	Seven	88	Female
15	Tade Row	Four	88	male
16	Gimmy	Four	88	male
25	Giff Tow	Seven	88	male
31	Marry Toeye	Four	88	Female
35	Rows Noup	Six	88	Female
28	Rojj Base	Seven	86	Female
2	Max Ruin	Three	85	male
8	Asruid	Five	85	male
27	Big Nose	Three	81	Female
23	Herod	Eight	79	male
26	Crelea	Seven	79	male
30	Reppy Red	Six	79	Female
7	My John Rob	Fifth	78	male
9	Tes Qry	Six	78	male

only showing top 20 rows

```
In [33]: df_pyspark.orderBy("mark").show()
```

id	name	class	mark	gender
----	------	-------	------	--------

Jupyter groupBy() Last Checkpoint: 10 hours ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (pykernel)

```
In [33]: df_pyspark.orderBy("mark").show()
```

id	name	class	mark	gender
19	Tinny	Nine	18	male
17	Tumyu	Six	54	male
3	Arnold	Three	55	male
6	Alex John	Four	55	male
10	Big John	Four	55	Female
22	Reggid	Seven	55	Female
29	Tess Played	Seven	55	male
4	Krish Star	Four	60	Female
5	John Mike	Four	60	Female
20	Jackly	Nine	65	Female
21	Babby John	Four	69	Female
34	Gain Toe	Seven	69	male
1	John Deo	Four	75	Female
18	Honny	Five	75	male
7	My John Rob	Fifth	78	male
9	Tes Qry	Six	78	male
24	Tiddy Now	Seven	78	male
23	Herod	Eight	79	male
26	Crelea	Seven	79	male
30	Reppy Red	Six	79	Female

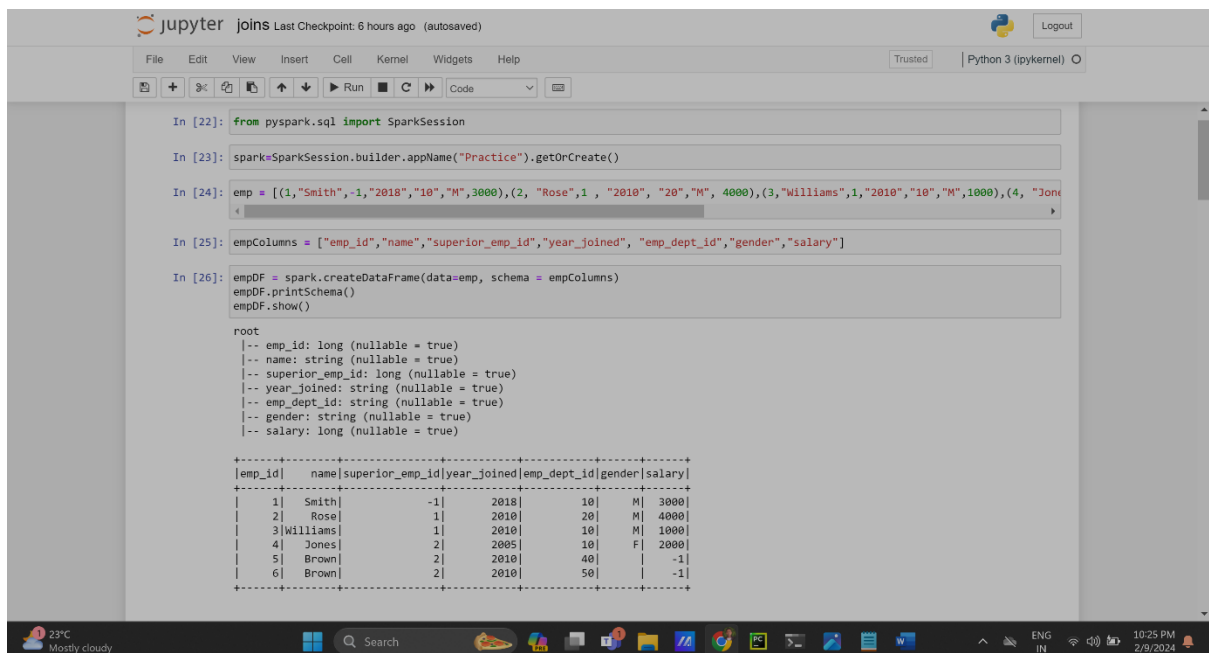
only showing top 20 rows

```
In [ ]:
```

join() using pyspark

PySpark Join is used to combine two DataFrames and by chaining these you can join multiple DataFrames; it supports all basic join type operations available in traditional SQL like INNER, LEFT OUTER, RIGHT OUTER, LEFT ANTI, LEFT SEMI, CROSS, SELF join.

Join Type	Description
Inner Join	Join records when key columns are matched, and dropped when they are not matched
Outer join	Returns all rows from both datasets, where Join expression doesn't match it returns null or respective columns
Left Join/ Left outer join	Returns all rows from left dataset regardless of match found on right dataset, when Join doesn't match – it assigns null for that record
Right Join/ Right outer join	Returns all rows from Right dataset regardless of match found on left dataset, when Join doesn't match – it assigns null for that record
Left Semi Join	Returns columns from the only left dataset for the matched records in the right dataset on join expression
Left Anti Join	Returns only columns from left dataset for non-matched records



```
In [22]: from pyspark.sql import SparkSession

In [23]: spark=SparkSession.builder.appName("Practice").getOrCreate()

In [24]: emp = [(1,"Smith",-1,"2018","10","M",3000),(2, "Rose",1, "2010", "20","M", 4000),(3,"Williams",1,"2010","10","M",1000),(4, "Jones",2,"2010","10","F",2000),(5,"Brown",2,"2010","40","M",4000),(6,"Brown",2,"2010","50","M",-1)]

In [25]: empColumns = ["emp_id","name","superior_emp_id","year_joined", "emp_dept_id","gender","salary"]

In [26]: empDF = spark.createDataFrame(data=emp, schema = empColumns)
empDF.printSchema()
empDF.show()

root
 |-- emp_id: long (nullable = true)
 |-- name: string (nullable = true)
 |-- superior_emp_id: long (nullable = true)
 |-- year_joined: string (nullable = true)
 |-- emp_dept_id: string (nullable = true)
 |-- gender: string (nullable = true)
 |-- salary: long (nullable = true)

+-----+-----+-----+-----+-----+-----+-----+
|emp_id|  name|superior_emp_id|year_joined|emp_dept_id|gender|salary|
+-----+-----+-----+-----+-----+-----+
|    1| Smith|          -1|    2018|         10|M|   3000|
|    2|  Rose|           1|    2010|         20|M|   4000|
|    3|Williams|          1|    2010|         10|M|   1000|
|    4|  Jones|           2|    2010|         10|F|   2000|
|    5| Brown|           2|    2010|         40|M|   4000|
|    6| Brown|           2|    2010|         50|M|    -1|
+-----+-----+-----+-----+-----+-----+

```

jupyter joins Last Checkpoint: 6 hours ago (autosaved) Python 3 (ipykernel) Logout

```
In [27]: dept = [{"Finance",10}, {"Marketing",20}, {"Sales",30}, {"IT",40}]
deptColumns = ["dept_name", "dept_id"]

In [28]: deptDF = spark.createDataFrame(data=dept, schema = deptColumns)

In [29]: deptDF.printSchema()
deptDF.show()

root
 |-- dept_name: string (nullable = true)
 |-- dept_id: long (nullable = true)

+-----+-----+
|dept_name|dept_id|
+-----+-----+
| Finance|    10|
|Marketing|    20|
|  Sales|    30|
|    IT|    40|
+-----+-----+

In [30]: #inner join

In [31]: empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"inner") .show()

+-----+-----+-----+-----+-----+-----+-----+-----+
|emp_id| name|superior_emp_id|year_joined|emp_dept_id|gender|salary|dept_name|dept_id|
+-----+-----+-----+-----+-----+-----+-----+-----+
| 1| Smith|      -1|    2018|      10|M| 3000| Finance|    10|
| 3|Williams|     1|    2010|      10|M| 1000| Finance|    10|
| 4| Jones|     2|    2005|      10|F| 2000| Finance|    10|
| 2| Rose|     1|    2010|      20|M| 4000|Marketing|    20|
| 5| Brown|     2|    2010|      40|M|    -1|    IT|    40|
```

jupyter joins Last Checkpoint: 6 hours ago (autosaved) Python 3 (ipykernel) Logout

```
In [30]: #inner join

In [31]: empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"inner") .show()

+-----+-----+-----+-----+-----+-----+-----+-----+
|emp_id| name|superior_emp_id|year_joined|emp_dept_id|gender|salary|dept_name|dept_id|
+-----+-----+-----+-----+-----+-----+-----+-----+
| 1| Smith|      -1|    2018|      10|M| 3000| Finance|    10|
| 3|Williams|     1|    2010|      10|M| 1000| Finance|    10|
| 4| Jones|     2|    2005|      10|F| 2000| Finance|    10|
| 2| Rose|     1|    2010|      20|M| 4000|Marketing|    20|
| 5| Brown|     2|    2010|      40|M|    -1|    IT|    40|
+-----+-----+-----+-----+-----+-----+-----+-----+

In [32]: #outer join

In [33]: empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"outer") .show()

+-----+-----+-----+-----+-----+-----+-----+-----+
|emp_id| name|superior_emp_id|year_joined|emp_dept_id|gender|salary|dept_name|dept_id|
+-----+-----+-----+-----+-----+-----+-----+-----+
| 1| Smith|      -1|    2018|      10|M| 3000| Finance|    10|
| 3|Williams|     1|    2010|      10|M| 1000| Finance|    10|
| 4| Jones|     2|    2005|      10|F| 2000| Finance|    10|
| 2| Rose|     1|    2010|      20|M| 4000|Marketing|    20|
| NULL| NULL|      NULL|      NULL|      NULL| NULL| NULL| Sales|    30|
| 5| Brown|     2|    2010|      40|M|    -1|    IT|    40|
| 6| Brown|     2|    2010|      50| |    -1| NULL| NULL|
+-----+-----+-----+-----+-----+-----+-----+-----+

In [34]: empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"full") .show()
```

Jupyter joins Last Checkpoint: 6 hours ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (ipykernel)

In [34]: `empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"full").show()`

emp_id	name	superior_emp_id	year_joined	emp_dept_id	gender	salary	dept_name	dept_id
1	Smith	-1	2018	10	M	3000	Finance	10
3	Williams	1	2010	10	M	1000	Finance	10
4	Jones	2	2005	10	F	2000	Finance	10
2	Rose	1	2010	20	M	4000	Marketing	20
NULL	NULL	NULL	NULL	NULL	NULL	NULL	Sales	30
5	Brown	2	2010	40		-1	IT	40
6	Brown	2	2010	50		-1	NULL	NULL

In [35]: `empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"fullouter").show()`

emp_id	name	superior_emp_id	year_joined	emp_dept_id	gender	salary	dept_name	dept_id
1	Smith	-1	2018	10	M	3000	Finance	10
3	Williams	1	2010	10	M	1000	Finance	10
4	Jones	2	2005	10	F	2000	Finance	10
2	Rose	1	2010	20	M	4000	Marketing	20
NULL	NULL	NULL	NULL	NULL	NULL	NULL	Sales	30
5	Brown	2	2010	40		-1	IT	40
6	Brown	2	2010	50		-1	NULL	NULL

In [36]: `#Left join`

In [37]: `empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"left").show()`

Jupyter joins Last Checkpoint: 6 hours ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (ipykernel)

In [36]: `#Left join`

In [37]: `empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"left").show()`

emp_id	name	superior_emp_id	year_joined	emp_dept_id	gender	salary	dept_name	dept_id
1	Smith	-1	2018	10	M	3000	Finance	10
2	Rose	1	2010	20	M	4000	Marketing	20
3	Williams	1	2010	10	M	1000	Finance	10
4	Jones	2	2005	10	F	2000	Finance	10
5	Brown	2	2010	40		-1	IT	40
6	Brown	2	2010	50		-1	NULL	NULL

In [38]: `empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"leftouter").show()`

emp_id	name	superior_emp_id	year_joined	emp_dept_id	gender	salary	dept_name	dept_id
1	Smith	-1	2018	10	M	3000	Finance	10
2	Rose	1	2010	20	M	4000	Marketing	20
3	Williams	1	2010	10	M	1000	Finance	10
4	Jones	2	2005	10	F	2000	Finance	10
5	Brown	2	2010	40		-1	IT	40
6	Brown	2	2010	50		-1	NULL	NULL

In [39]: `#right join`

In [40]: `empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"right").show()`

Jupyter joins Last Checkpoint: 6 hours ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (ipykernel)

In [39]: `#right join`

In [40]: `empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"right").show()`

emp_id	name	superior_emp_id	year_joined	emp_dept_id	gender	salary	dept_name	dept_id
4	Jones	2	2005	10	F	2000	Finance	10
3	Williams	1	2010	10	M	1000	Finance	10
1	Smith	-1	2018	10	M	3000	Finance	10
2	Rose	1	2010	20	M	4000	Marketing	20
NULL	NULL	NULL	NULL	NULL	NULL	NULL	Sales	30
5	Brown	2	2010	40		-1	IT	40

In [41]: `empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"rightouter").show()`

emp_id	name	superior_emp_id	year_joined	emp_dept_id	gender	salary	dept_name	dept_id
4	Jones	2	2005	10	F	2000	Finance	10
3	Williams	1	2010	10	M	1000	Finance	10
1	Smith	-1	2018	10	M	3000	Finance	10
2	Rose	1	2010	20	M	4000	Marketing	20
NULL	NULL	NULL	NULL	NULL	NULL	NULL	Sales	30
5	Brown	2	2010	40		-1	IT	40

In [42]: `#Leftsemi Join`

In [43]: `empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"leftsemi").show()`

Jupyter joins Last Checkpoint: 6 hours ago (autosaved) Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (ipykernel)

In [42]: `#Leftsemi Join`

In [43]: `empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"leftsemi").show()`

emp_id	name	superior_emp_id	year_joined	emp_dept_id	gender	salary
1	Smith	-1	2018	10	M	3000
3	Williams	1	2010	10	M	1000
4	Jones	2	2005	10	F	2000
2	Rose	1	2010	20	M	4000
5	Brown	2	2010	40		-1

In [44]: `#Left Anti Join`

In [45]: `empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"leftanti").show()`

emp_id	name	superior_emp_id	year_joined	emp_dept_id	gender	salary
6	Brown	2	2010	50		-1

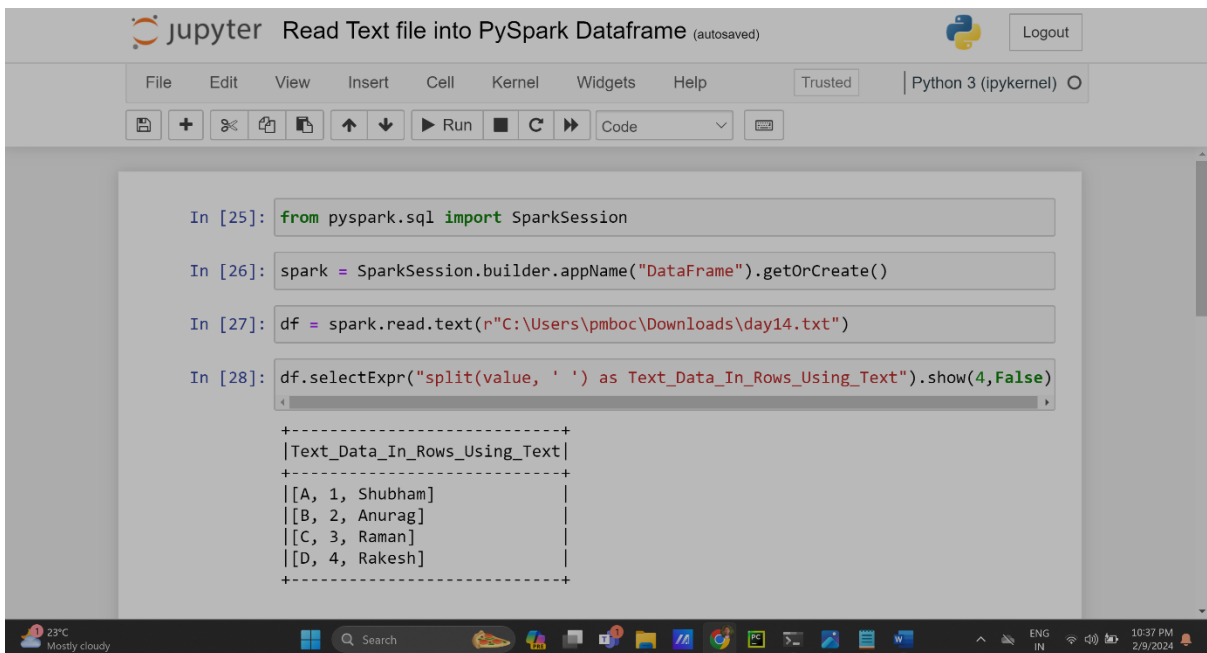
In []:

Read Text file into PySpark Dataframe

- Using `spark.read.text()`
- Using `spark.read.csv()`
- Using `spark.read.format().load()`

❖ Method 1: Using `spark.read.text()`

It is used to load text files into DataFrame whose schema starts with a string column. Each line in the text file is a new row in the resulting DataFrame. Using this method we can also read multiple files at a time.



```
In [25]: from pyspark.sql import SparkSession

In [26]: spark = SparkSession.builder.appName("DataFrame").getOrCreate()

In [27]: df = spark.read.text(r"C:\Users\pmboc\Downloads\day14.txt")

In [28]: df.selectExpr("split(value, ' ') as Text_Data_In_Rows_Using_Text").show(4,False)

+-----+
|Text_Data_In_Rows_Using_Text|
+-----+
|[A, 1, Shubham]|
|[B, 2, Anurag]|
|[C, 3, Raman]|
|[D, 4, Rakesh]|
+-----+
```

❖ Method 2: Using spark.read.csv()

It is used to load text files into DataFrame. Using this method we will go through the input once to determine the input schema if inferSchema is enabled. To avoid going through the entire data once, disable inferSchema option or specify the schema explicitly using the schema.

```
In [29]: df = spark.read.csv(r"C:\Users\pmboc\Downloads\day14.txt")

In [30]: df.selectExpr("split(_c0, ' ') as Text_Data_In_Rows_Using_CSV").show(4,False)

+-----+
|Text_Data_In_Rows_Using_CSV|
+-----+
|[A, 1, Shubham]|
|[B, 2, Anurag]|
|[C, 3, Raman]|
|[D, 4, Rakesh]|
+-----+
```

❖ Method 3: Using spark.read.format()

It is used to load text files into DataFrame. The .format() specifies the input data source format as “text”. The .load() loads data from a data source and returns DataFrame.

```
In [31]: df = spark.read.format("text").load(r"C:\Users\pmboc\Downloads\day14.txt")

In [32]: df.selectExpr("split(value, ' ') as Text_Data_In_Rows_Using_format_load").show(4)

+-----+
|Text_Data_In_Rows_Using_format_load|
+-----+
|[A, 1, Shubham]|
|[B, 2, Anurag]|
|[C, 3, Raman]|
|[D, 4, Rakesh]|
+-----+
```